CONTROL-EF

Htoo Lwin (120832)
Fathima Shafana (121985)

Background

- The recent COVID-19 pandemic has forced many educational institutions at all levels to conduct lectures online in lieu of physical classes.
- Most such online lectures are recorded by the institution for student review or administrative purposes and are hence are usually opened to student access.
- Students can use these videos to either review course material or learn about new concepts from courses not directly related to their curriculum.

Problem Statement

- Navigating through the lecture content can prove to be time consuming and difficult without any prior processing of the videos.
- The main difficulties are **twofold**:
  1. Students can take some time manually going through video lists to find the lectures of topics they would like to learn (if the videos are publicly accessible in the first place).
  2. Students can also have a hard time searching through a particular video for a time when the lecturer is talking about the topic they want to learn about, especially if that student is reviewing the lecture.

Proposed Solution

- An application that helps users to get the exact timestamp of the keyword searched for.
- An architecture that is easily scalable and has **good read performance.**

Aims & Objectives

1. To provide a platform for students to easily access video lectures in an educational institution.
2. To provide students with the ability to quickly search for videos or timestamps using different facets such as tags, keywords.

Quality Attribute Analysis

Availability	**H**	The platform's main purpose is to provide aid in the learning process, so it should be available as much as possible. It should not fail the student when he needs it.
Performance	**H**	When a student is learning, it is important not to lose focus. So, the platform's search function should work as quickly as possible.
Portability	**L**	As the majority of quality studying is done in front of a desk, the software architecture should be designed for desktop computers (Windows PC, Mac OS).
Security	**M**	The main function - searching - only retrieves data, which is not dangerous. But uploading and transcript editing should be done only by authorized personnel. The software should provide that.
Scalability	**H**	As this platform is basically a database, any kind of expansion should be seamless and unnoticed by the user..
Testability	**M**	The functionality is quite simple and input methods are limited. The upload formats (video and audio) are also limited. The only functionality that should be tested thoroughly is the transcript generation.

Architecture Design and Implementation

Architectural Challenges and Other Discussion

- Database Performance
  - **Transcript search performance should be consistent regardless of data volume**.
    - Has to search every row of transcript table
  - Our servers won't handle video/transcript processing so lightweight! Only handles **Transcript search**.

Architectural Challenges and Other Discussion

- Transcript-Timestamp Mapping in Database
  - Inspired by YouTube’s caption format
- Integration
  - YouTube API authentication and integration
    - API keys/OAuth clients (register our app as client?)
  - Upload and get transcripts

Integration Challenges

- Transcript retrieval
  - No clear documentation on public APIs
    - Different public APIs and different parameters
    - Unreliable
  - Secure APIs need OAuth 2.0 clients for authentication
  - Official documentation only supports secure APIs.
- Upload Video
  - Same issue with authentication

Integration Challenges - Public API troubles

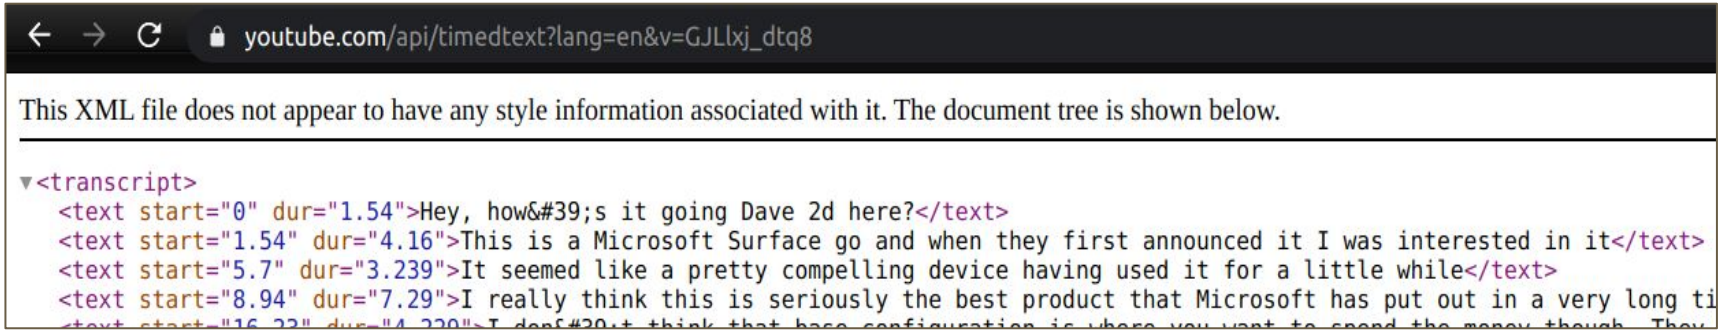

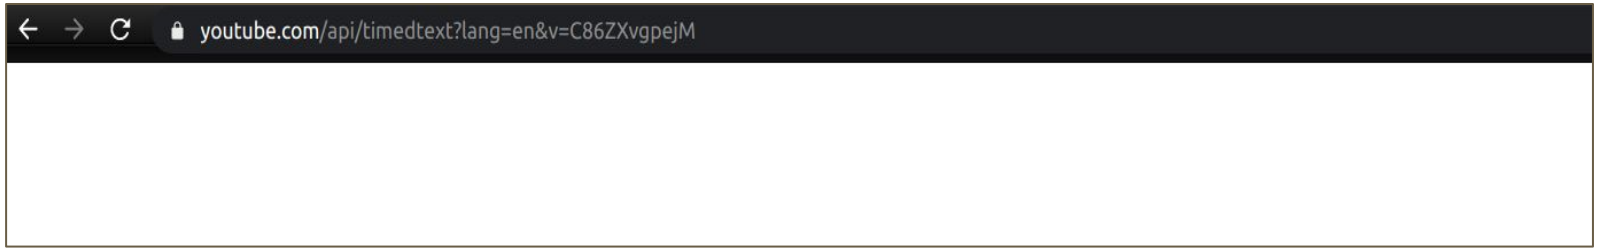

Architectural Challenges and Other Discussion

- Scalability (to support more users)
  - Our architecture should be able to easily scale and store more transcripts and videos
  - Our servers won’t handle video/transcript processing so lightweight! Only handles Transcript search.

Proposed Architectures - Single Server

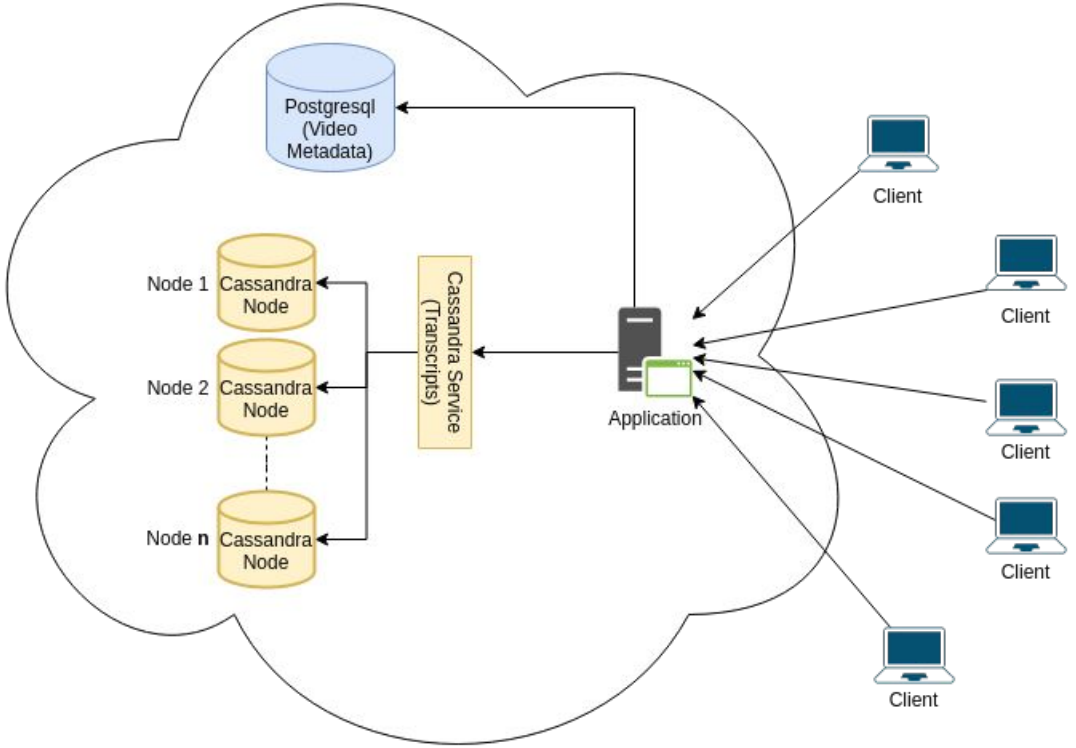

Proposed Architectures - Multi Server

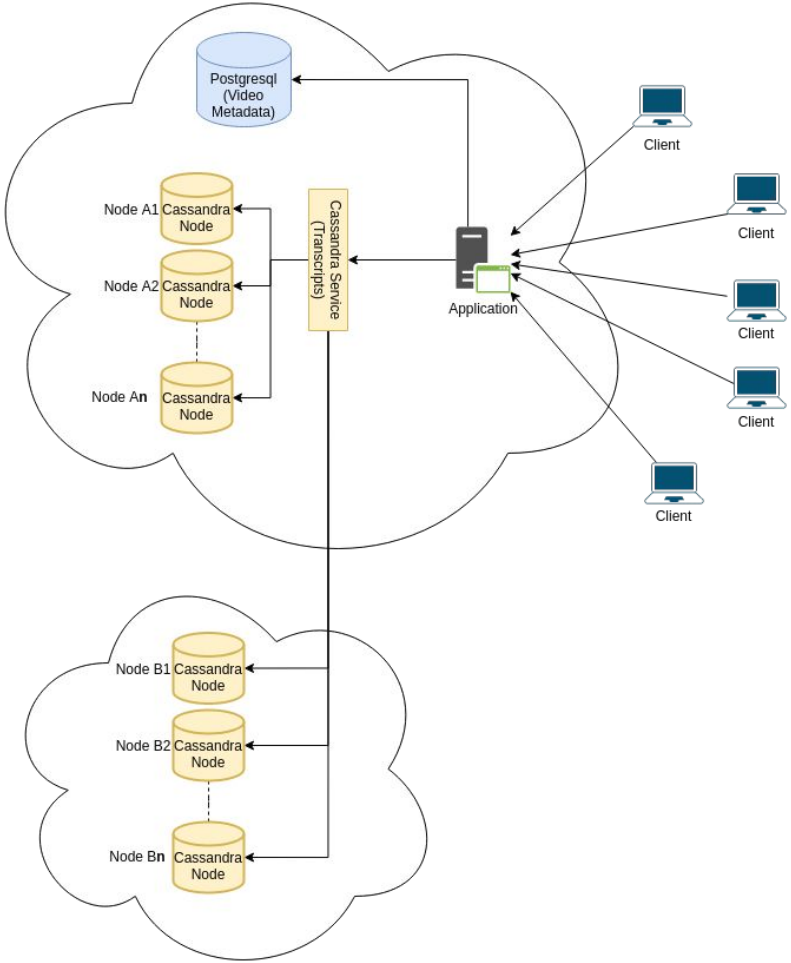

Architecture Patterns

1. Model-view-controller (MVC)
2. Service-oriented Architecture
3. Peer-to-Peer(?)

- The Control-EF application will utilize the **model-view-controller** (MVC) architecture pattern in main using the MVC framework Spring Boot.
- It will also incorporate some best practices from **service-oriented architecture** (SOA). Students will consume the videos from YouTube and transcripts from our database.

Architectural Design - Deployment Diagram

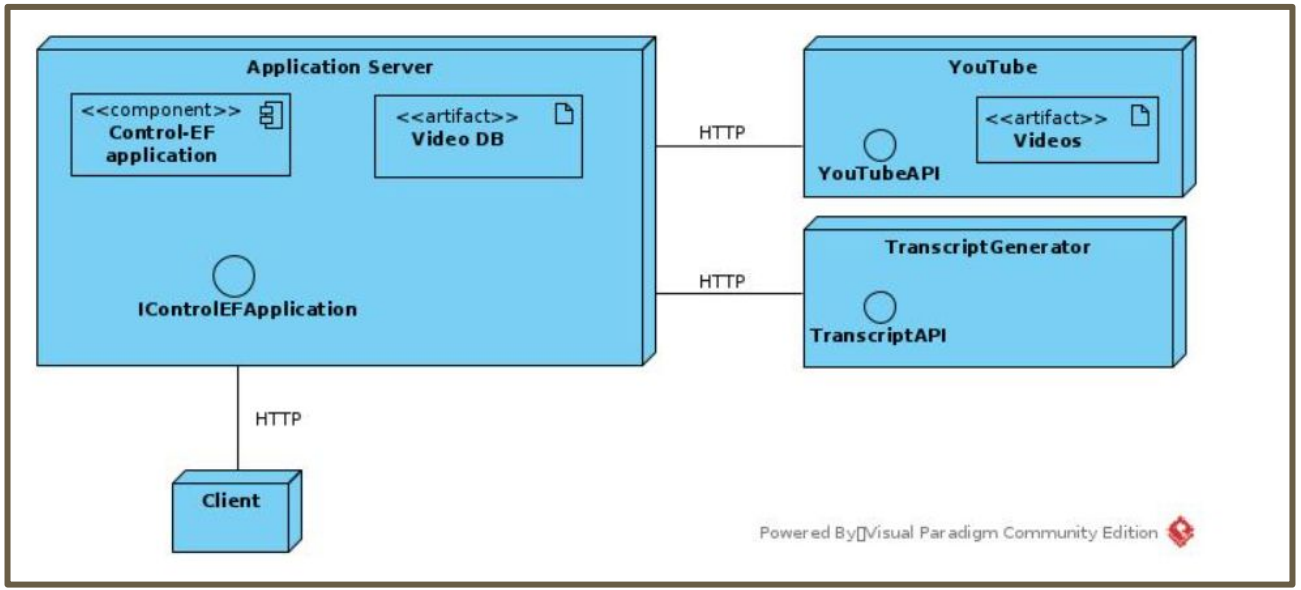

Class Diagram

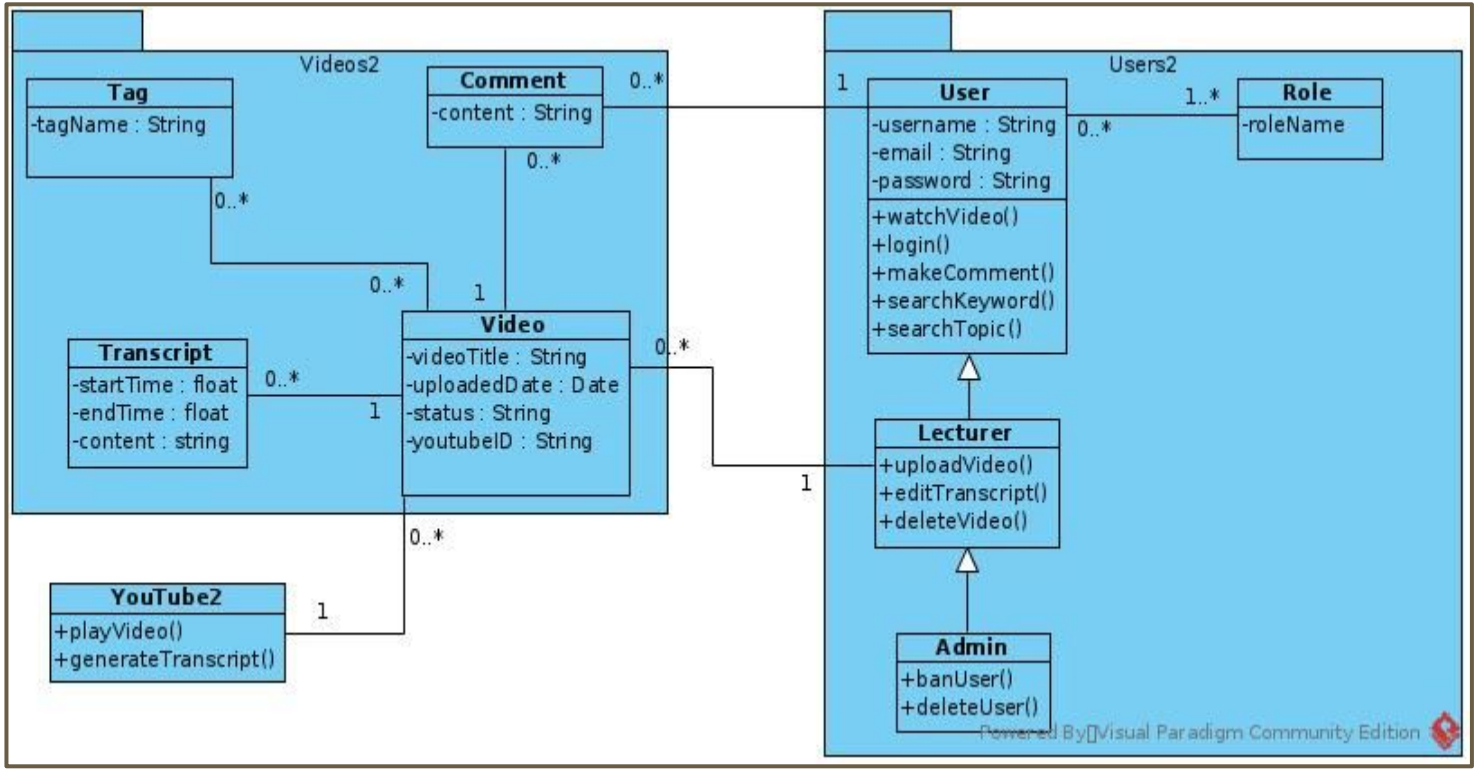

Tools and Technologies

Database: Postgres (use JDBC Postgres Driver)

Language: Java 8

Application Framework: Spring Boot

Testing Tools: Apache JMeter & Cassandra-Stress

Video-Transcript Mapping

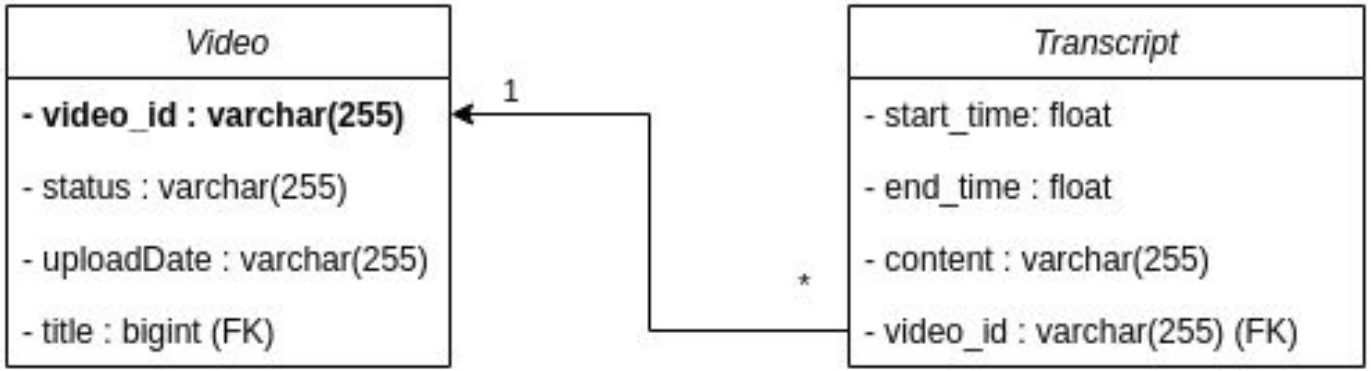

Transcripts (Ours vs. Youtube)

id	content	end_time	start_time	video_id
830	How do you observe something you can't see?	18260	15260	c8re1U9rCo4
831	This is the basic question of somebody who's interested	21260	18260	c8re1U9rCo4
832	in finding and studying black holes.	23260	21260	c8re1U9rCo4
833	Because black holes are objects	25260	23260	c8re1U9rCo4
834	whose pull of gravity is so intense	28260	25260	c8re1U9rCo4
835	that nothing can escape it, not even light,	30260	28260	c8re1U9rCo4

This XML file does not appear to have any style information associated with it. The document tree is shown below.


```
<transcript>
  <text start="0" dur="1.54">Hey, how&#39;s it going Dave 2d here?</text>
  <text start="1.54" dur="4.16">This is a Microsoft Surface go and when they first announced it I was interested in it</text>
  <text start="5.7" dur="3.239">It seemed like a pretty compelling device having used it for a little while</text>
  <text start="8.94" dur="7.29">I really think this is seriously the best product that Microsoft has put out in a very long time this thing starts at $400</text>
  <text start="16.23" dur="4.229">I don&#39;t think that base configuration is where you want to spend the money though. They have a mid tier one</text>
  <text start="21.16" dur="5.029">550 quite a bit more but you&#39;re getting double the RAM double the storage but significantly faster storage</text>
  <text start="26.26" dur="4.49">That is the model that I think most people should pick up if you can afford that price bump</text>
  <text start="30.75" dur="3.299">so this unit here, is that mid tier model the</text>
  <text start="34.78" dur="2">$550 unit and I</text>
  <text start="37.42" dur="5.209">Really like it. Ok, let&#39;s go around. This thing build quality is great. It&#39;s a surface product</text>
  <text start="42.629" dur="3.21">It has a magnesium enclosure fit and finish on this is really well done</text>
  <text start="45.84" dur="0.64">the</text>
```

Database Stress Testing

Testing Objectives

1. To test our proposed database architecture for consistent database performance for searching transcripts.

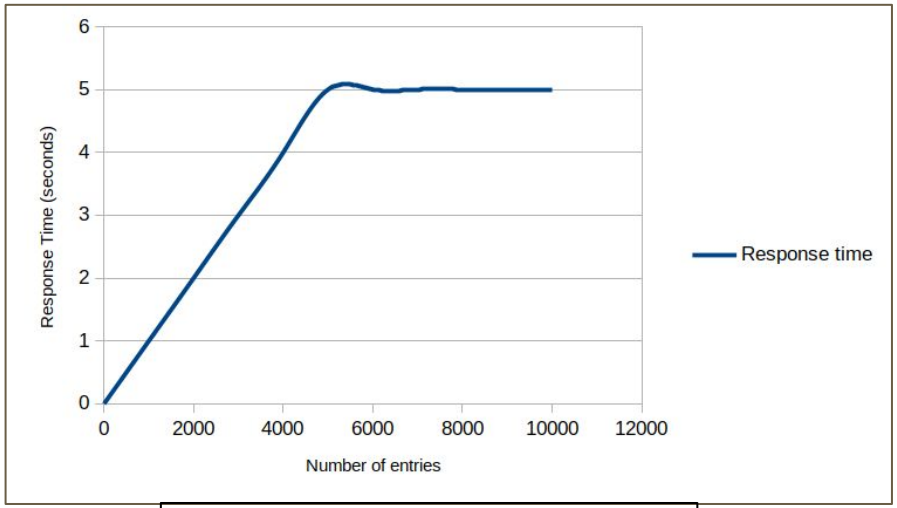

Our expected (and desired) result.

Test Methodology

- Compare Postgres search performance vs. Cassandra search performance
- Simulate keyword searching in real scenario with this query:
  - `select * from transcripts where content like '%a%'`
- Test at different amount of data records
  - 2k, 5k, 10k, 25k, 50k, 100k, 300k and 500k records
- **Find the mean latency in the response time of 200 read queries in milliseconds and compare**
  - **Evaluation metric - mean response time**
- Run tests on both Postgres and Cassandra

Test Environment

- The tests are run on the Guppy server available at CSIM here at AIT.
- Why Guppy?
  - To maintain the same and standard test environment.
  - Local Machines are of different specifications.
  - Control other variable except the Database
  - Databases themselves are located in Docker containers on Guppy
- For cassandra, two nodes were used = two containers

Test Architecture - Postgres

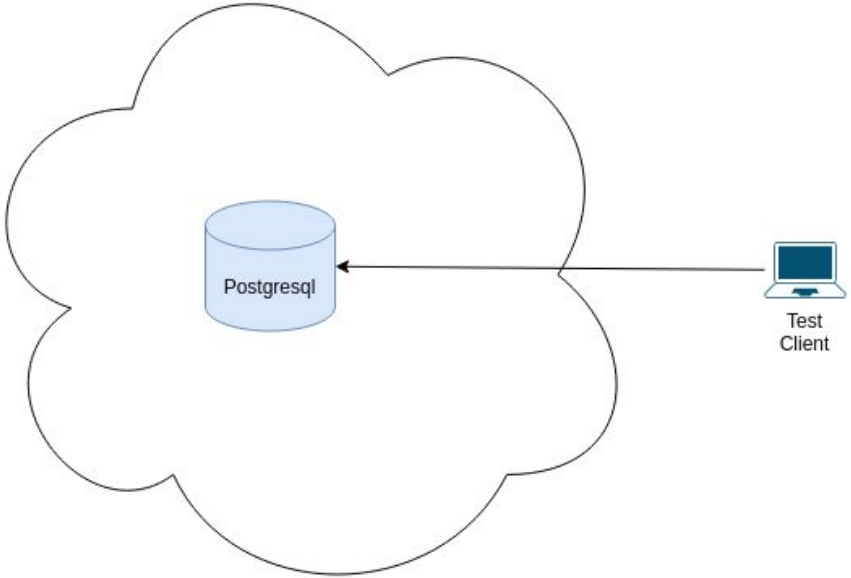

Test Architecture - Cassandra

Two nodes

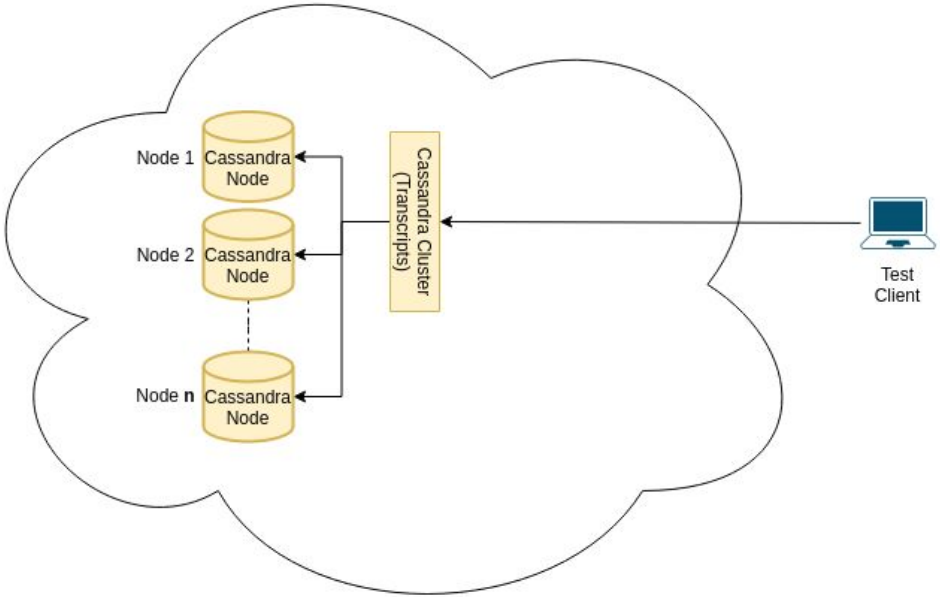

Test Architecture - Cassandra contd.

```
st120832@guppy:~$ docker exec -it control-cassandra2 nodetool status
Datacenter: datacenter1
=======================
Status=Up/Down
|/ State=Normal/Leaving/Joining/Moving
--  Address     Load       Tokens       Owns (effective)  Host ID                               Rack
UN  172.19.0.3  16.28 MiB  256          48.0%             18caa8bd-0f90-4f56-b6da-1f057d354ef7  rack1
UN  172.19.0.2  18.82 MiB  256          52.0%             e833e23a-c5b1-4ec4-a888-1dcc379b7f50  rack1
```

Two nodes with 500k records in the same datacenter/cluster

Test Architecture - Cassandra contd.

See this link in the Github Repository for instructions on how to set up this cluster

Test Tools

1. **Apache JMeter 5.4.1**
   - JMeter is a widely used tool utilized for analyzing and measuring the performance of a variety of services.
   - In the case of this study, it is used for the particular case of database stress testing.
   - JMeter natively supports most relational databases through JDBC driver plugins and this is the case for PostgreSQL as well.
2. **Cassandra-Stress Tool**
   - Cassandra is not supported natively by Apache JMeter
   - Plugins such as this are sorely out of date
   - Cassandra Stress tool was found to be more suitable for our testing purposes and is officially supported by the Apache foundation.

Test Scenario 1: PostgreSQL

Phase 1: Prepare test data sets for Stress Testing

- According to the Test Plan, a series of tables with records ranging from 2000, 5000, 10000, 25000, 50000, 100000, 300000 and 500000 were built in PostreSQL database as transcript1, transcript2, transcript3, transcript4, transcript5, transcript6, transcript7 and transcript8 respectively.

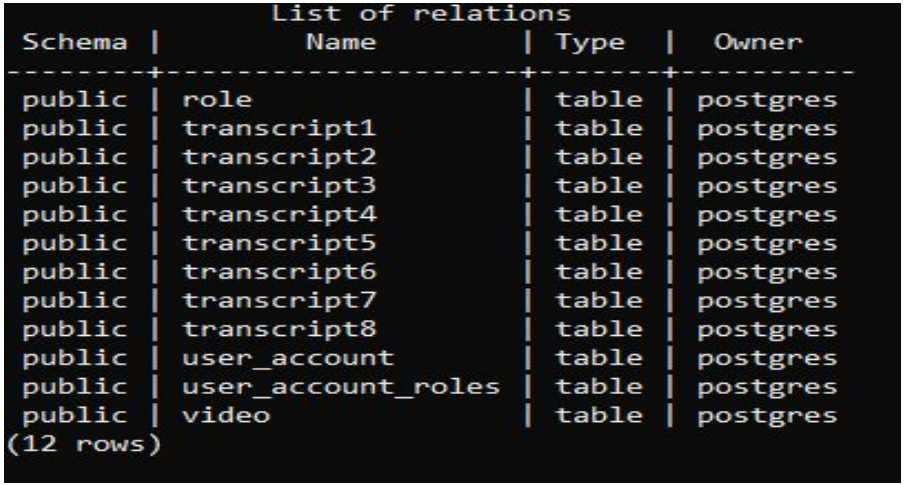

```
                List of relations
 Schema |        Name         | Type  |  Owner
--------+---------------------+-------+----------
 public | role                | table | postgres
 public | transcript1         | table | postgres
 public | transcript2         | table | postgres
 public | transcript3         | table | postgres
 public | transcript4         | table | postgres
 public | transcript5         | table | postgres
 public | transcript6         | table | postgres
 public | transcript7         | table | postgres
 public | transcript8         | table | postgres
 public | user_account        | table | postgres
 public | user_account_roles  | table | postgres
 public | video               | table | postgres
(12 rows)
```

Test Scenario 1: PostgreSQL contd.

Phase 2: Stress Testing with Apache JMeter 5.4.1

- Configure JDBC connection for controlefdb

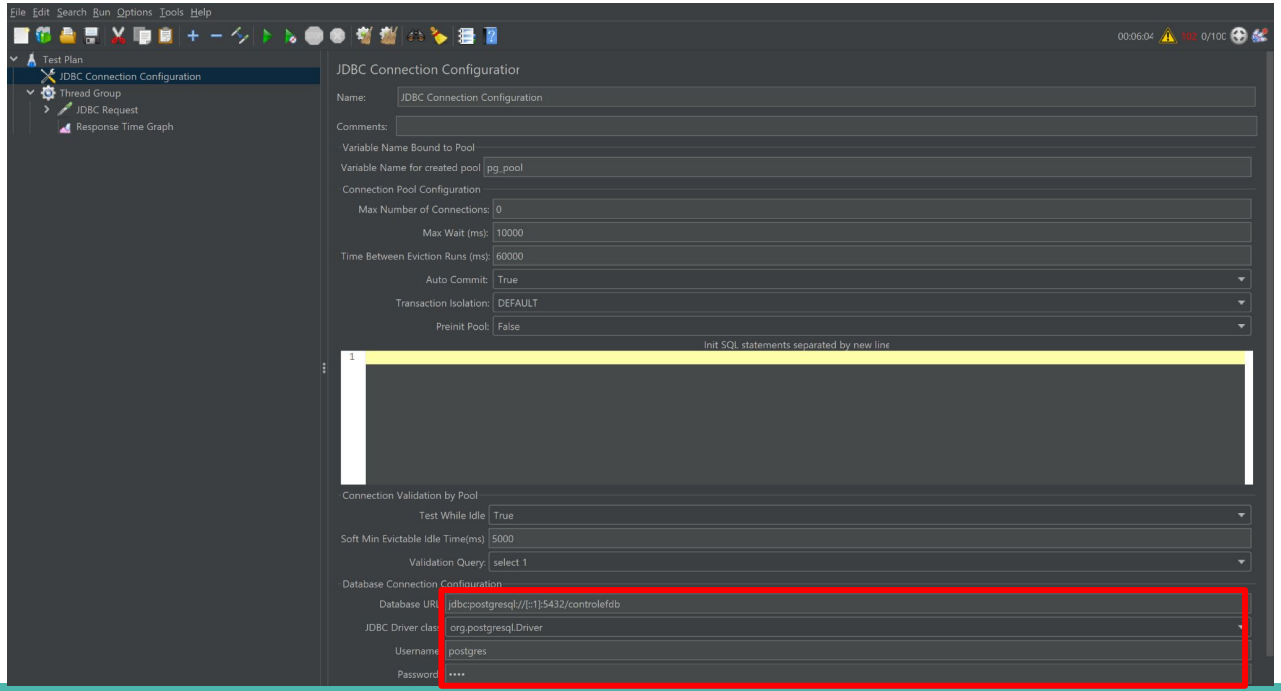

Test Scenario 1: PostgreSQL contd.

- Create Thread Group

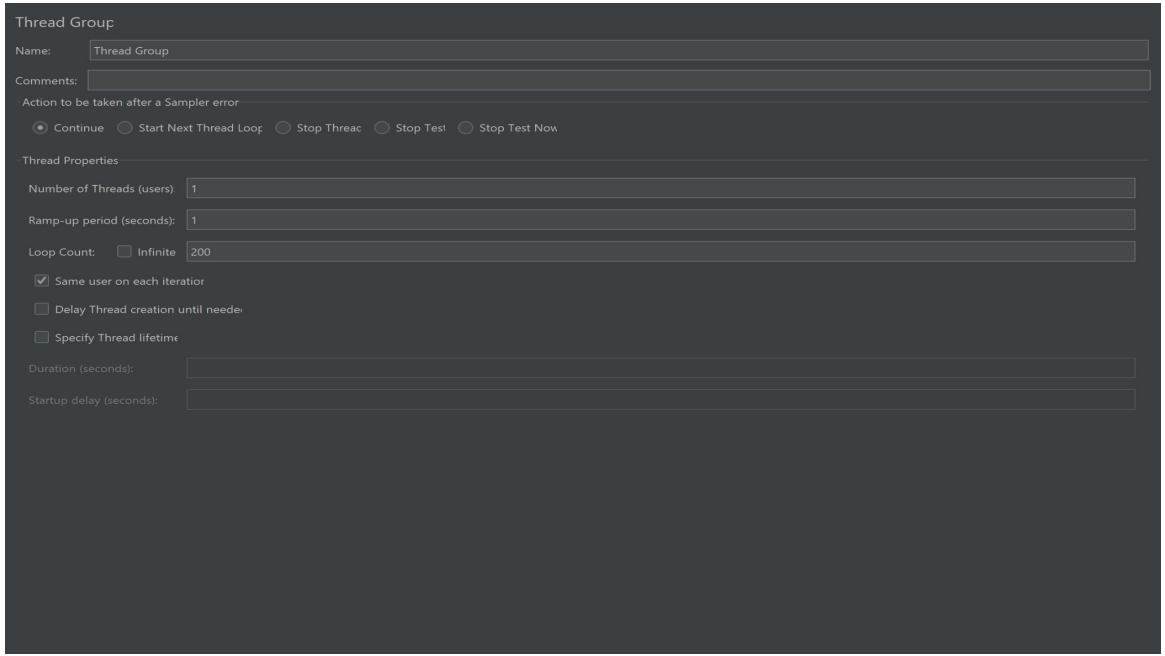

Test Scenario 1: PostgreSQL contd.

- Create Thread Group

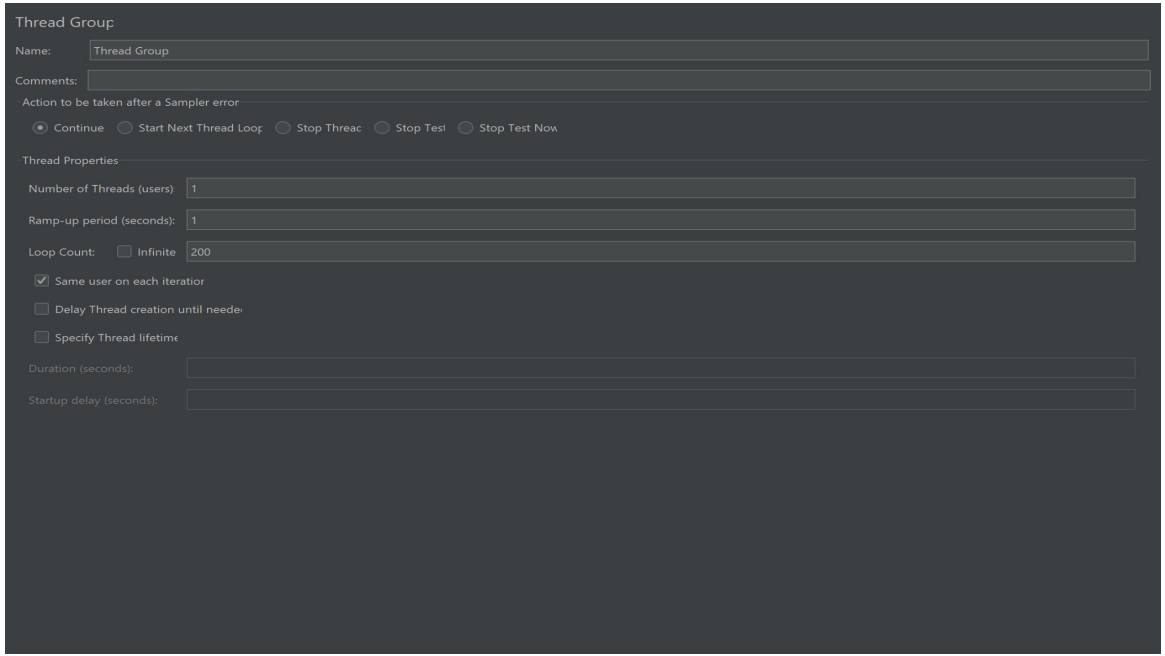

Test Scenario 1: PostgreSQL contd.

- Add JDBC Request as Sampler

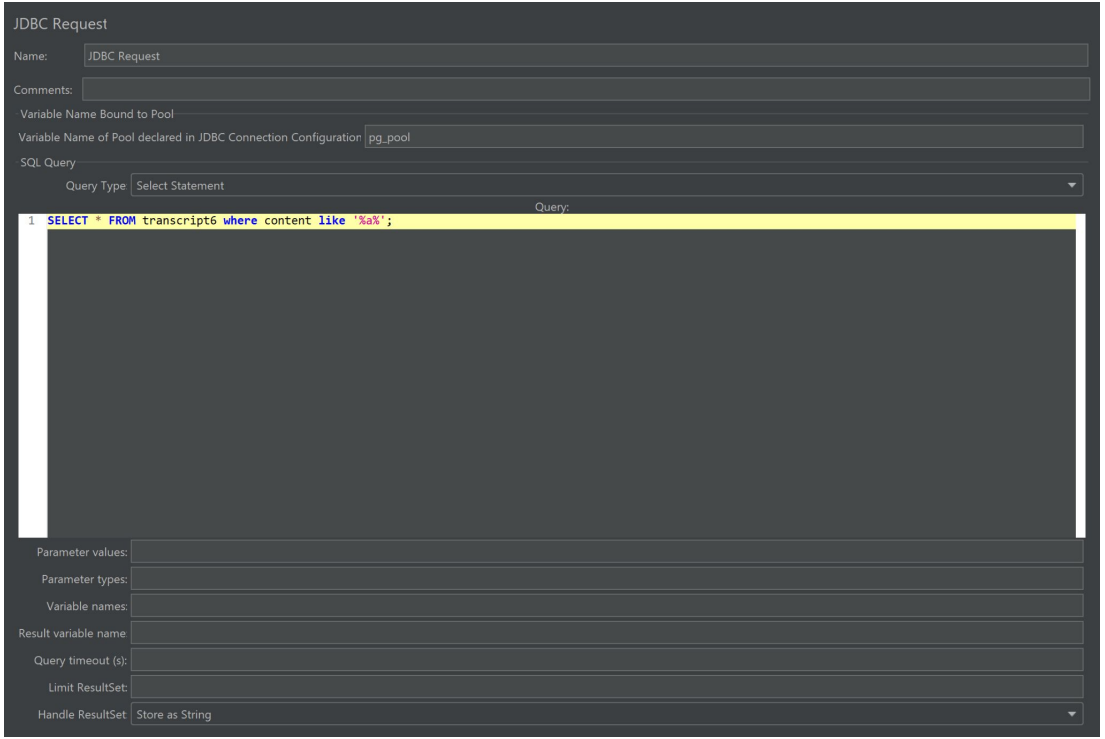

Test Scenario 1: PostgreSQL contd.

- Add listeners to view your results
  - View Results Tree
  - Graph Results
  - Aggregate Graph

Result: No. of Transcripts - 2 000

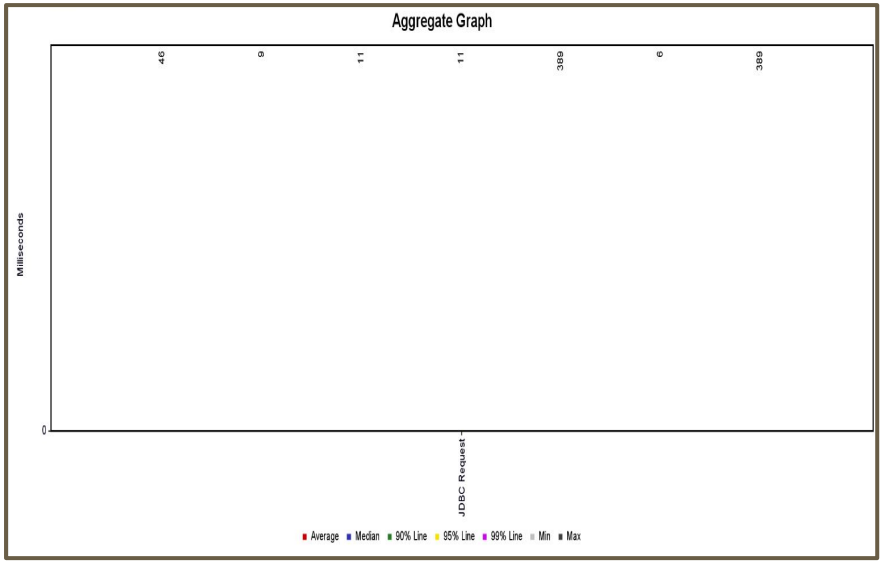

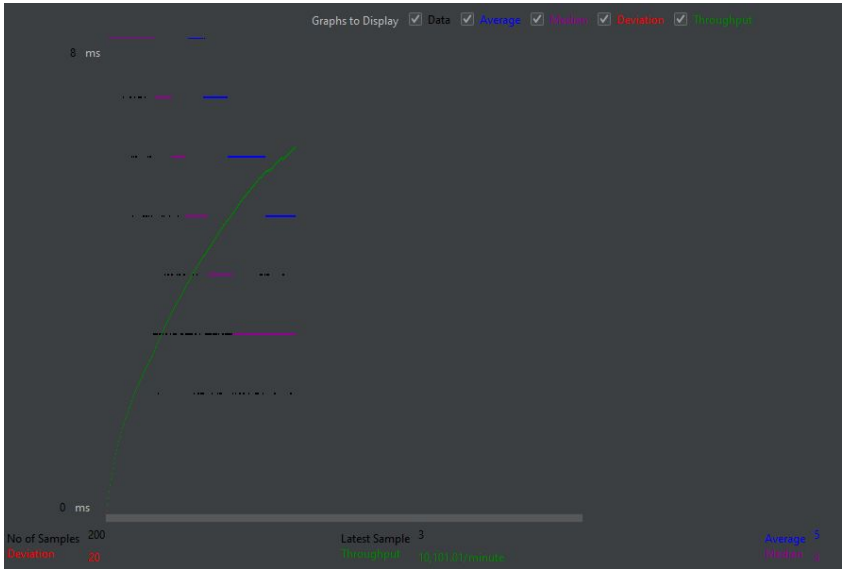

Average Latency = 5 ms

Result: No. of Transcripts - 5 000

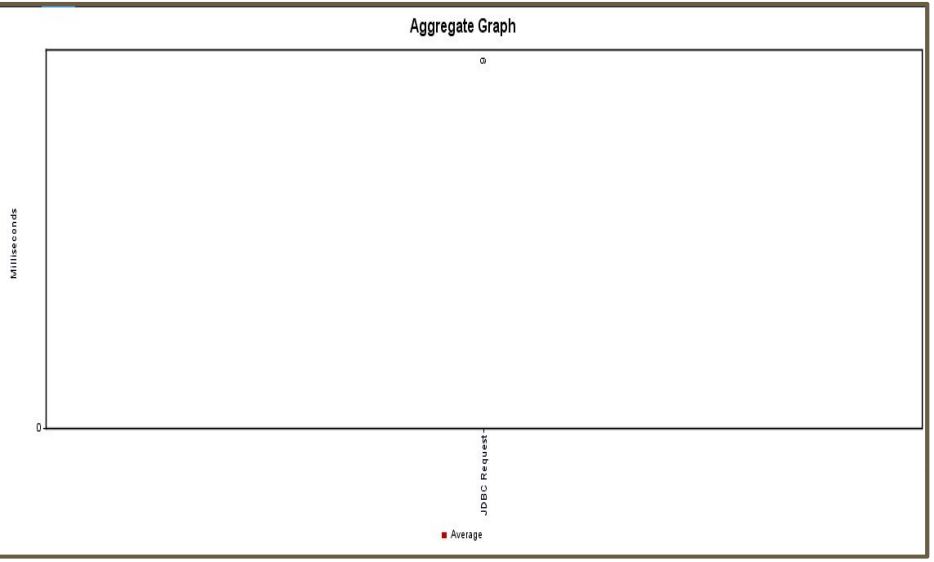

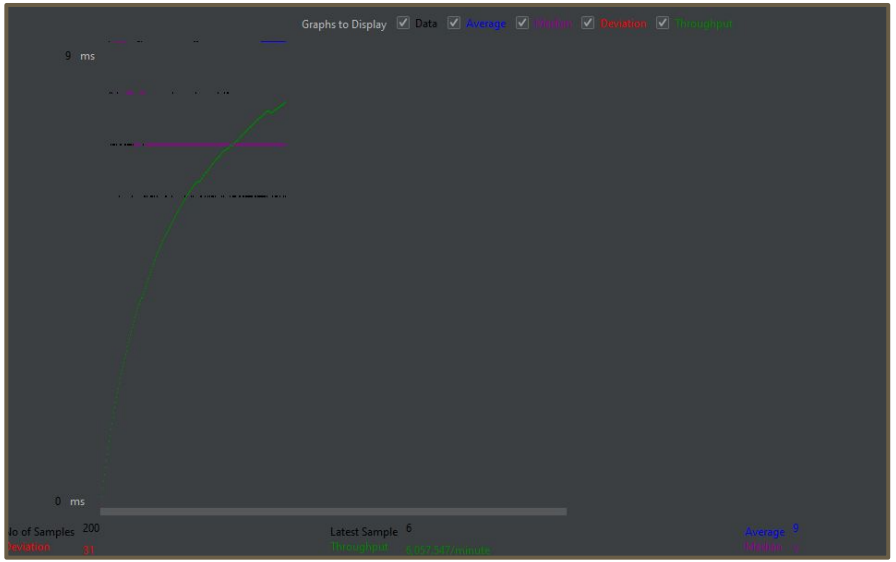

Average Latency = 9 ms

Result: No. of Transcripts - 10 000

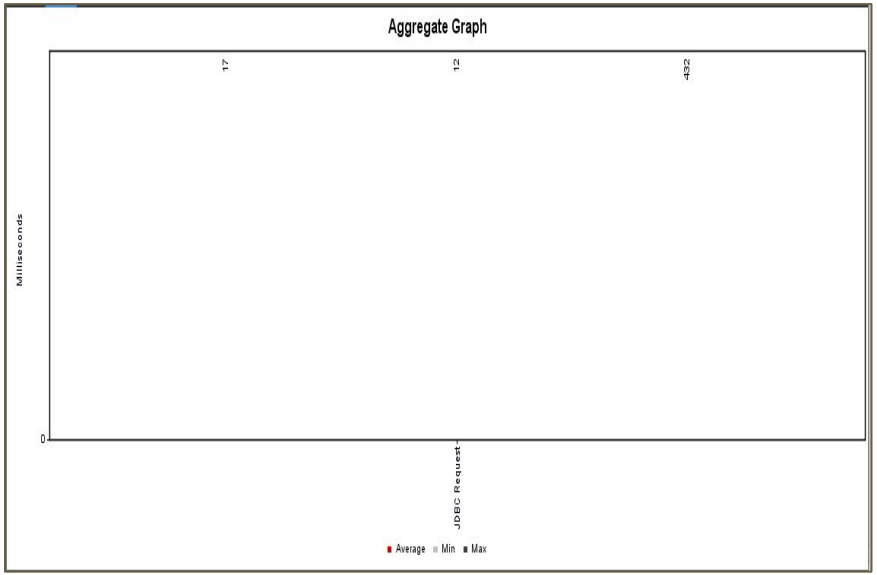

Average Latency = 17 ms

Result: No. of Transcripts - 25 000

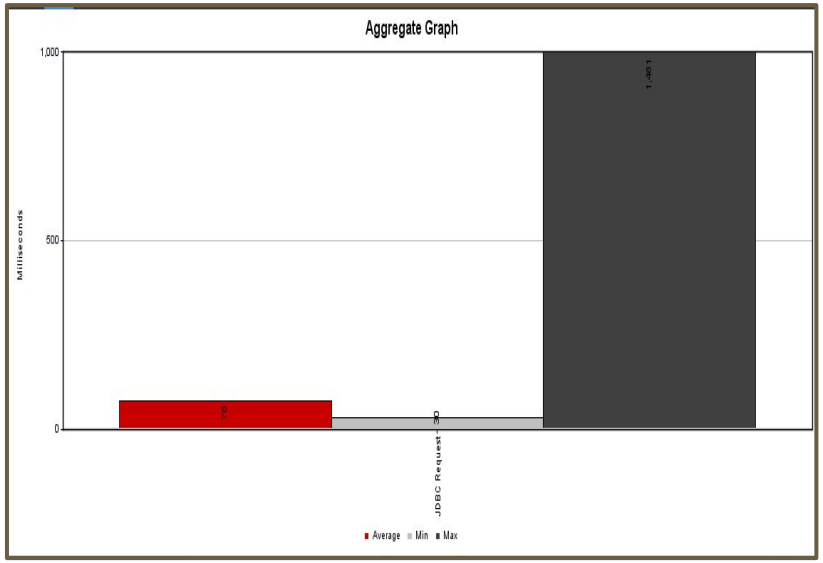

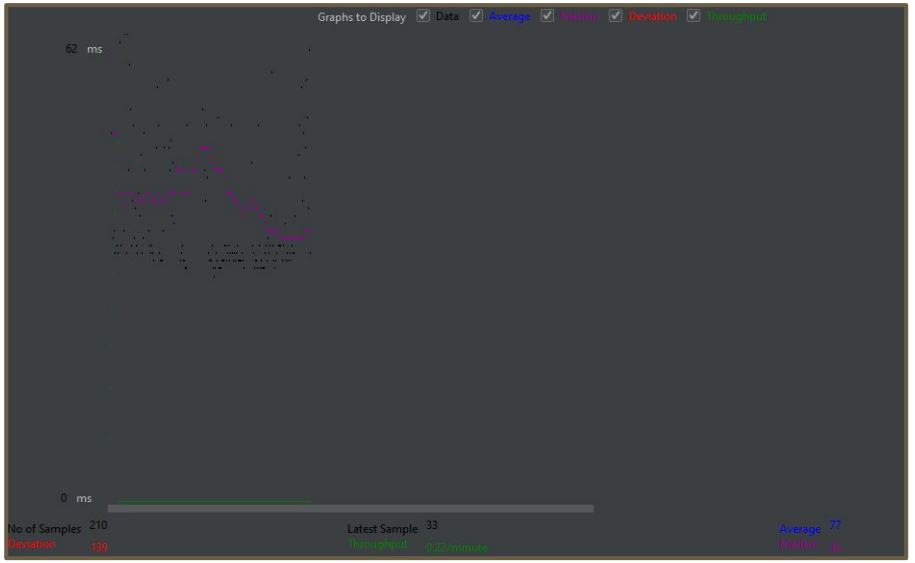

Average Latency = 77 ms

Result: No. of Transcripts - 50 000

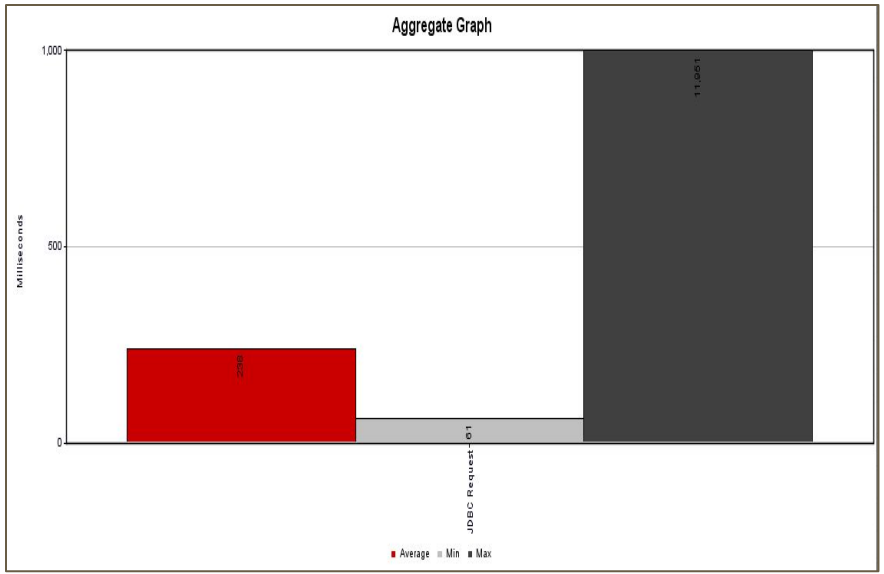

Average Latency = 238 ms

Result: No. of Transcripts - 100 000

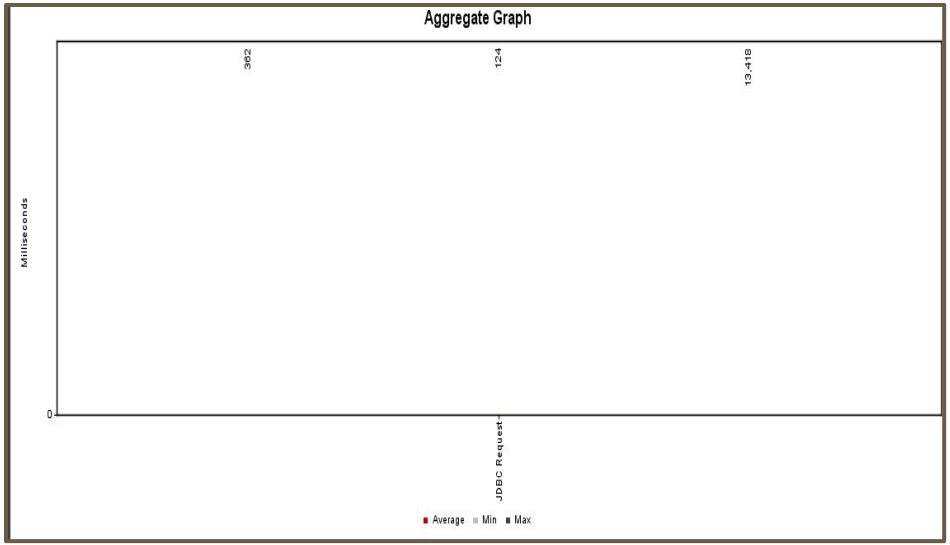

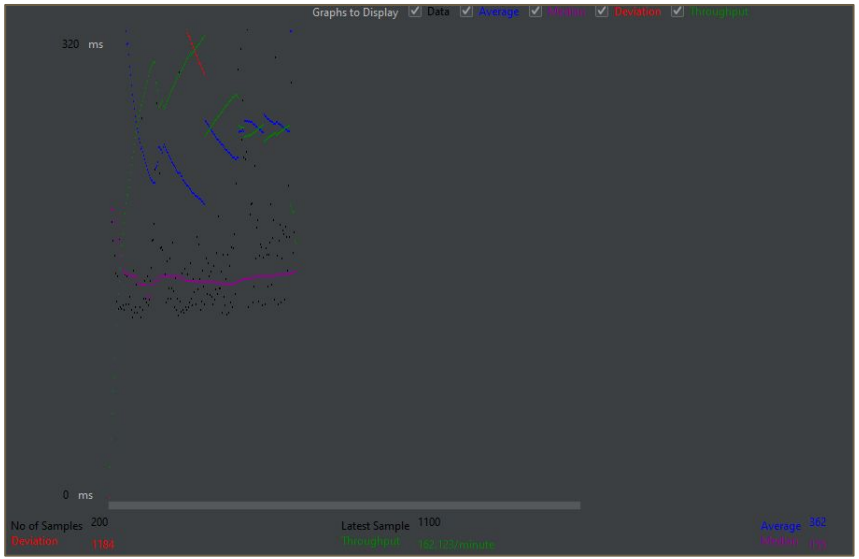

Average Latency = 362 ms

Result: No. of Transcripts - 300 000

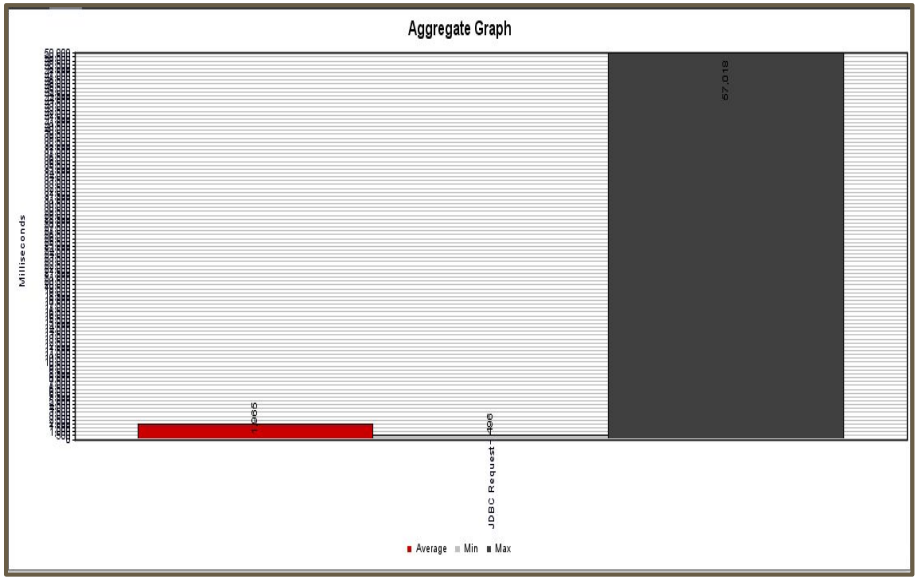

Average Latency = 1965 ms

Result: No. of Transcripts - 500 000

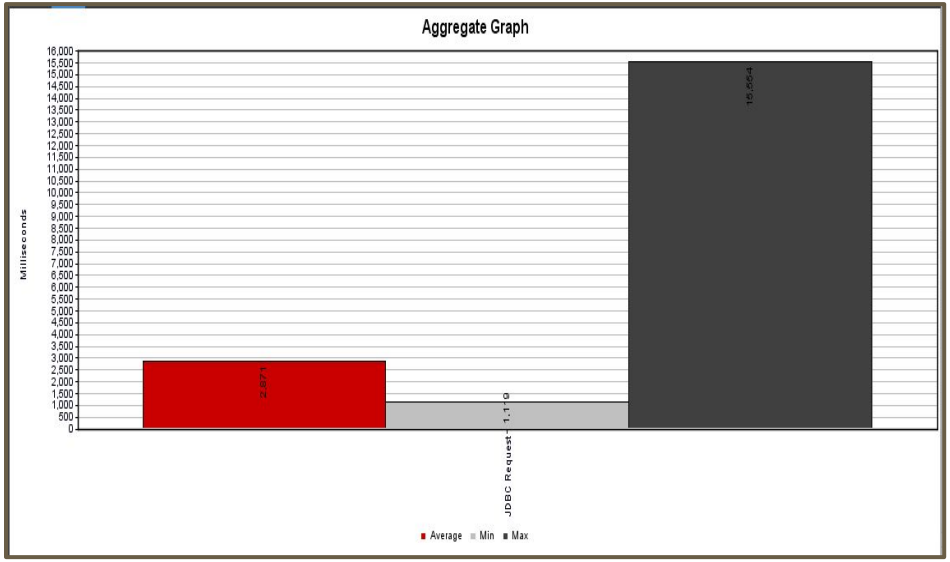

Average Latency = 2871 ms

Test Scenario 2: Cassandra

- Use YAML file for config to populate and read data

```yaml
keyspace: transcript

table: transcripts_by_content

columnspec:
  - name: content
    size: uniform(5..50)
  - name: video_id
    size: fixed(11)

insert:
  # How many partition to insert per batch
  partitions: fixed(1)
  # How many rows to update per partition
  select: fixed(1)/500
  # UNLOGGED or LOGGED batch for insert
  batchtype: UNLOGGED

queries:
  read1:
    cql: select * from transcripts_by_content where content like '%a%'
    fields: samerow
```

Test Scenario 2: Cassandra contd.

- Run stress tool to populate the database.

```
adam@adam-Prestige-14-A10SC:~/apache-cassandra-3.11.10/tools/bin$ ./cassandra-stress user profile=controlef.yml n=100000 cl=ONE ops\(insert=1\) -rate threads=1 -graph file=t
est.html title=test revision=test1 -node localhost,localhost:4127,localhost:4128
****************** Stress Settings ******************
Command:
  Type: user
  Count: 100,000
  No Warmup: false
  Consistency Level: ONE
  Target Uncertainty: not applicable
  Command Ratios: {insert=1.0}
  Command Clustering Distribution: clustering=gaussian(1..10)
  Profile File: controlef.yml
Rate:
  Auto: false
```

- `n` here refers to batches - basically the number of times the insert will be run.

Test Scenario 2: Cassandra contd.

- Run stress tool to read the database.
- Generates an html with a graph of the results

```
adam@adam-Prestige-14-A10SC:~/apache-cassandra-3.11.10/tools/bin$ ./cassandra-stress user profile=controlef.yml n=200 cl=ONE ops\(read1=1\) -rate threads=1 -graph file=test.
html title=test revision=test1 -node localhost
****************** Stress Settings ******************
Command:
  Type: user
  Count: 200
  No Warmup: false
  Consistency Level: ONE
  Target Uncertainty: not applicable
  Command Ratios: {read1=1.0}
  Command Clustering Distribution: clustering=gaussian(1..10)
  Profile File: controlef.yml
Rate:
  Auto: false
```

Test Scenario 2: Cassandra contd. - Read Results

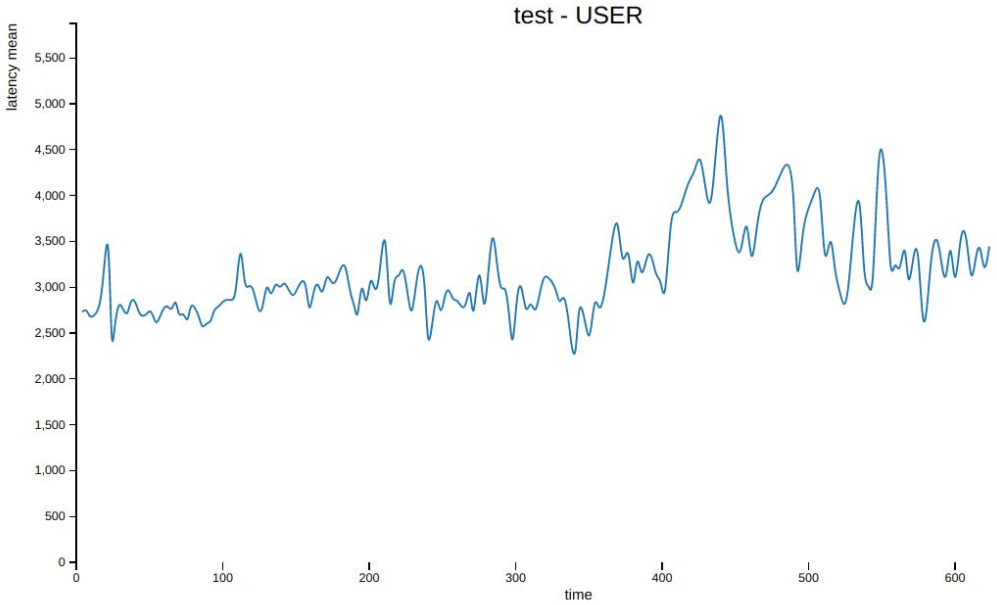

Choose metric: latency mean
Choose operation: USER
Data smoothing: 1
Show aggregates
Zoom: reset
x min 0 x max 685.74
y min 0 y max 5875.54

To hide/show a dataset click on the associated colored box

test1

```
----------------------------------------
op rate                   : 0 op/s [read1: 0 op/s]
partition rate            : 0 pk/s [read1: 0 pk/s]
row rate                  : 10,817 row/s [read1: 10,870 row/s]
latency mean              : 3114.0 ms [read1: 3,114.0 ms]
latency median            : 2975.9 ms [read1: 2,975.9 ms]
latency 95th percentile   : 4135.6 ms [read1: 4,135.6 ms]
latency 99th percentile   : 4521.5 ms [read1: 4,521.5 ms]
latency 99.9th percentile : 5343.5 ms [read1: 5,343.5 ms]
latency max               : 5343.5 ms [read1: 5,343.5 ms]
total gc count            : 0
total gc memory           : 0.000 KiB
total gc time             : 0.0 seconds
avg gc time               : NaN ms
stddev gc time            : 0.0 ms
Total operation time      : 00:10:23
cmd: user profile=controlef.yml n=200 cl=ONE ops(read1=1) -rate threads=1 -graph file=test.html title=test revision=test1 -node localhost -pop seq=1..100000
```

Test Scenario 2: Cassandra contd.

- Repeat for other data record levels (2k, 5k, 10k, **50k (shown below)** etc.)

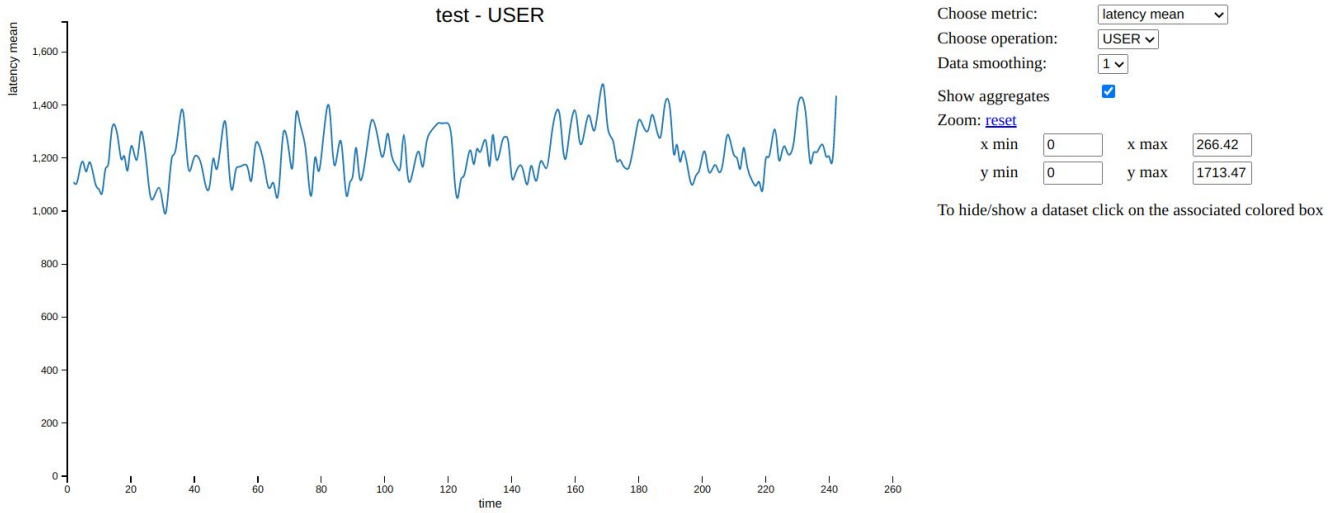

Test Scenario 2 - Cassandra contd.

See this link in the Github Repository for instructions on how to run cassandra-stress

Final Result - Comparing Postgres and Cassandra

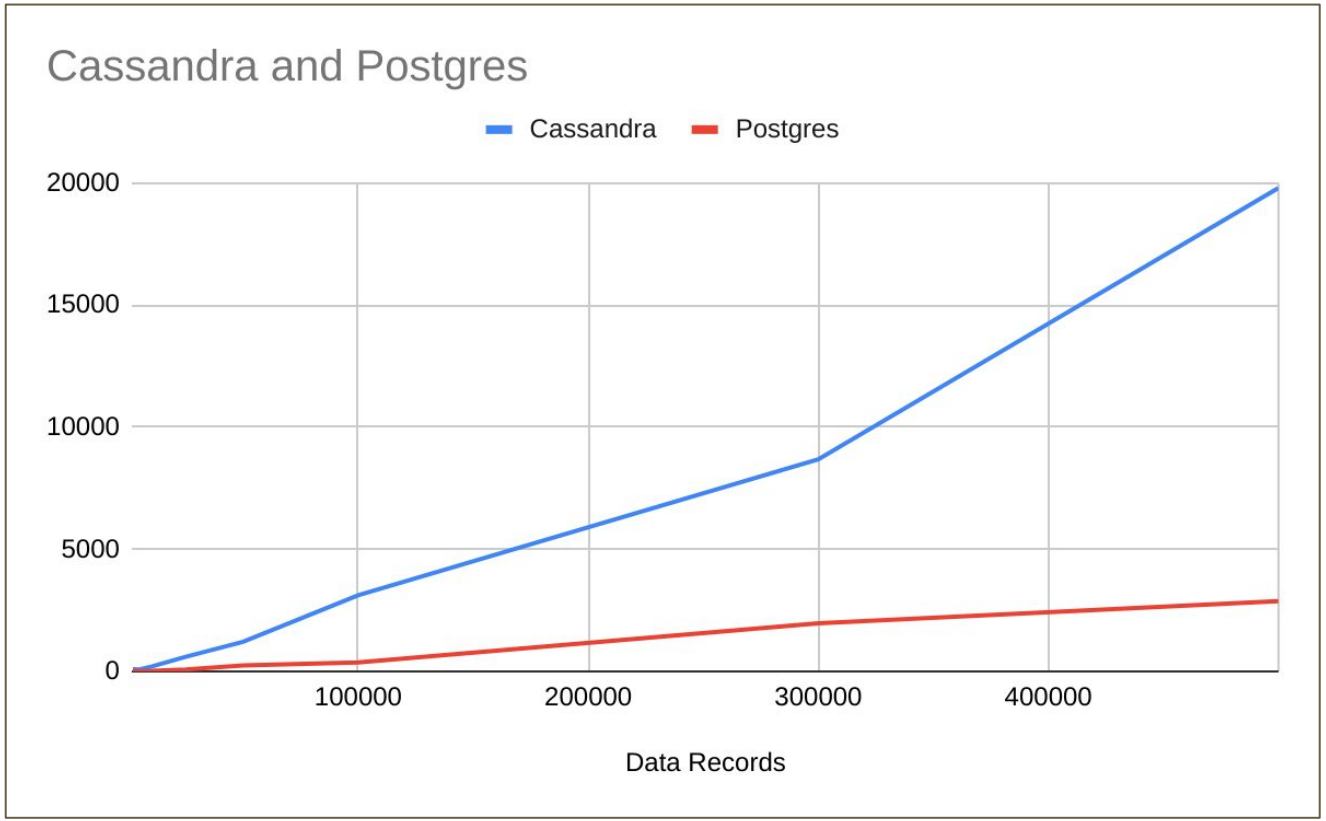

Final Result - Comparing Postgres and Cassandra

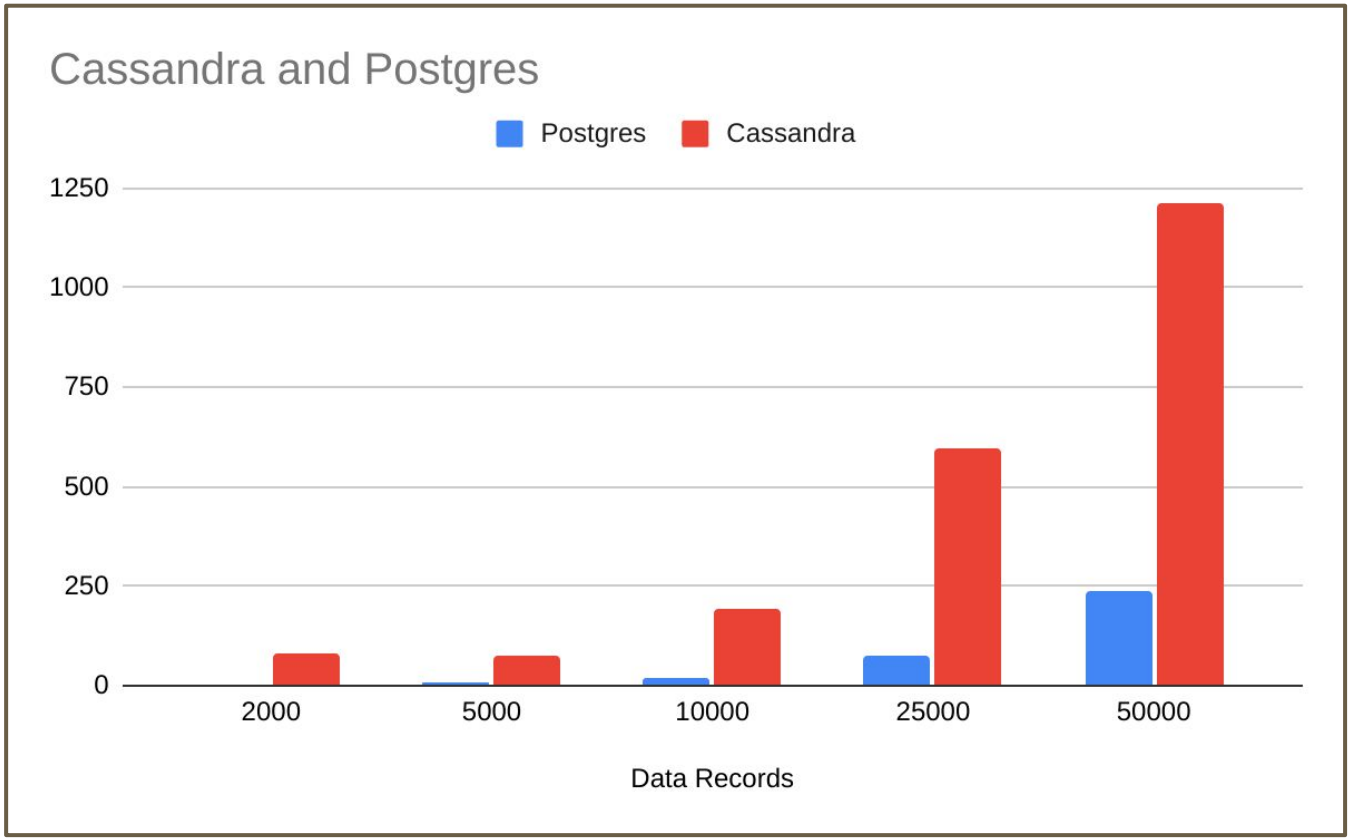

Final Result - Comparing Postgres and Cassandra

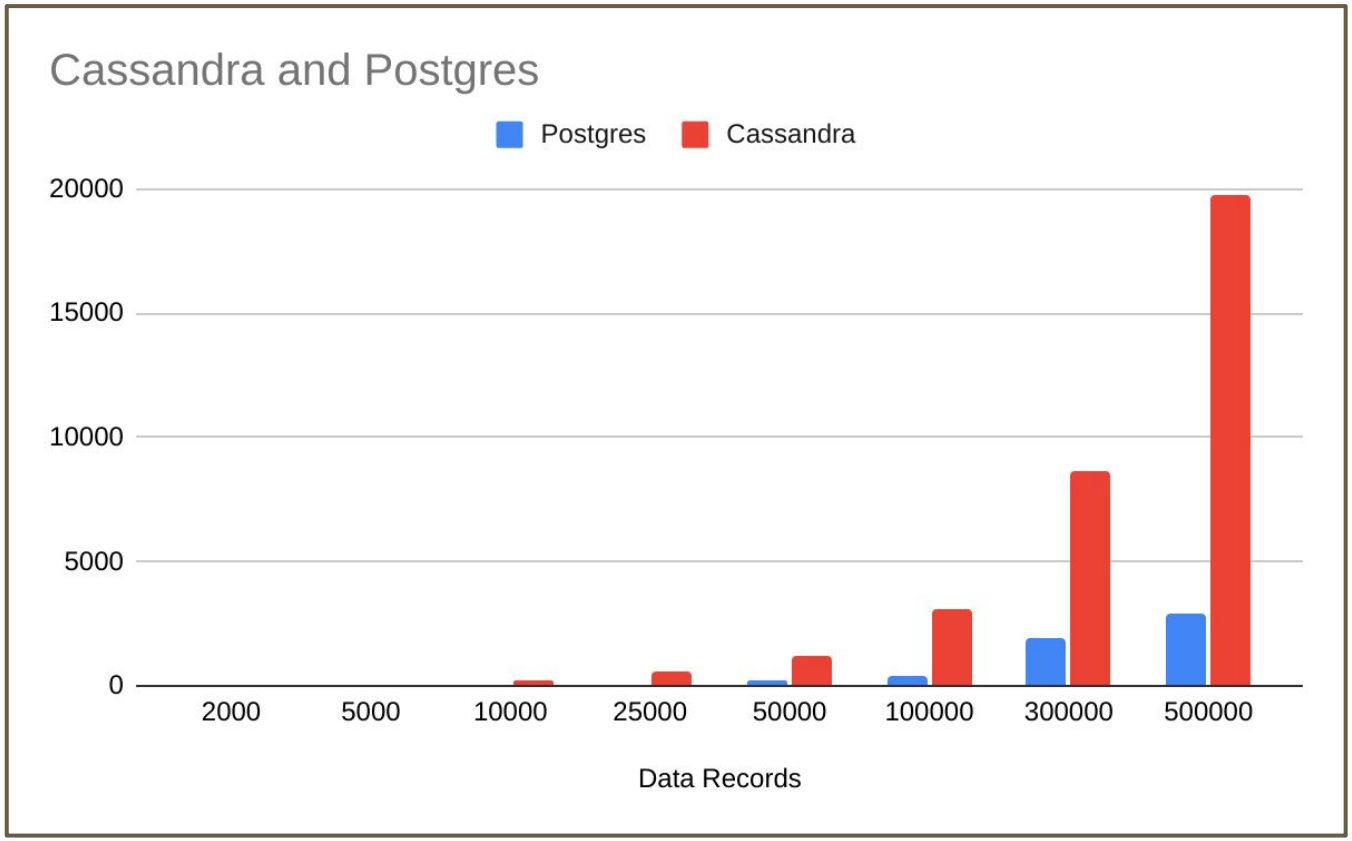

Relation to other findings

- Workload C (Read-only)
- https://adataanalyst.com/data-analysis-resources/a-comparison-between-cassandra-and-mysql/

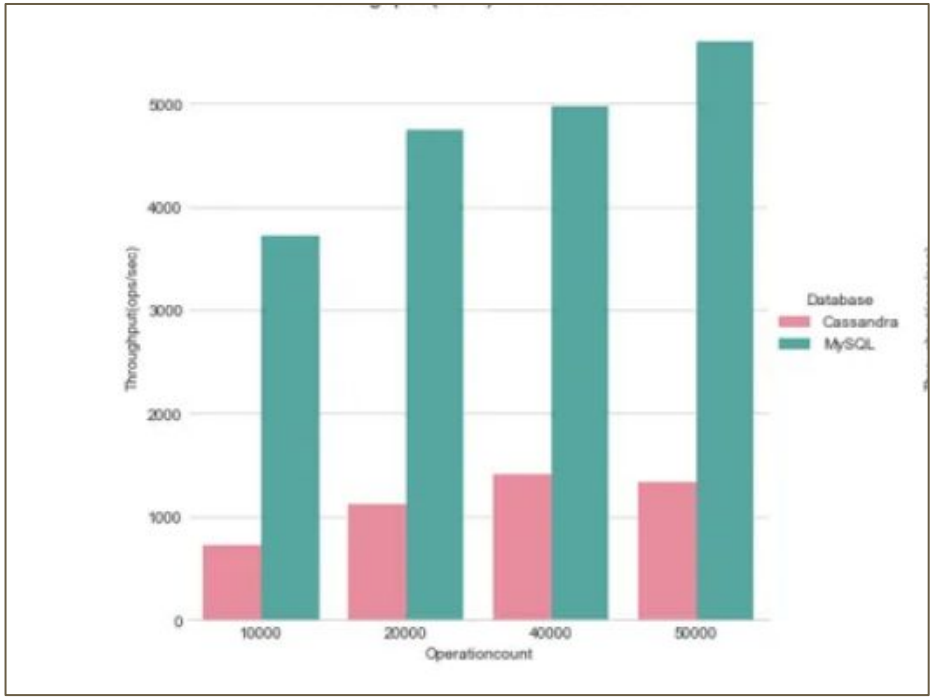

Relation to other findings (contd.)

- Mahmood, K. (2016). Performance Comparison of NOSQL Database Cassandra and SQL Server for Large Databases. *Journal of Independent Studies and Research (JISR)*, *14*(2).
- But questionable?

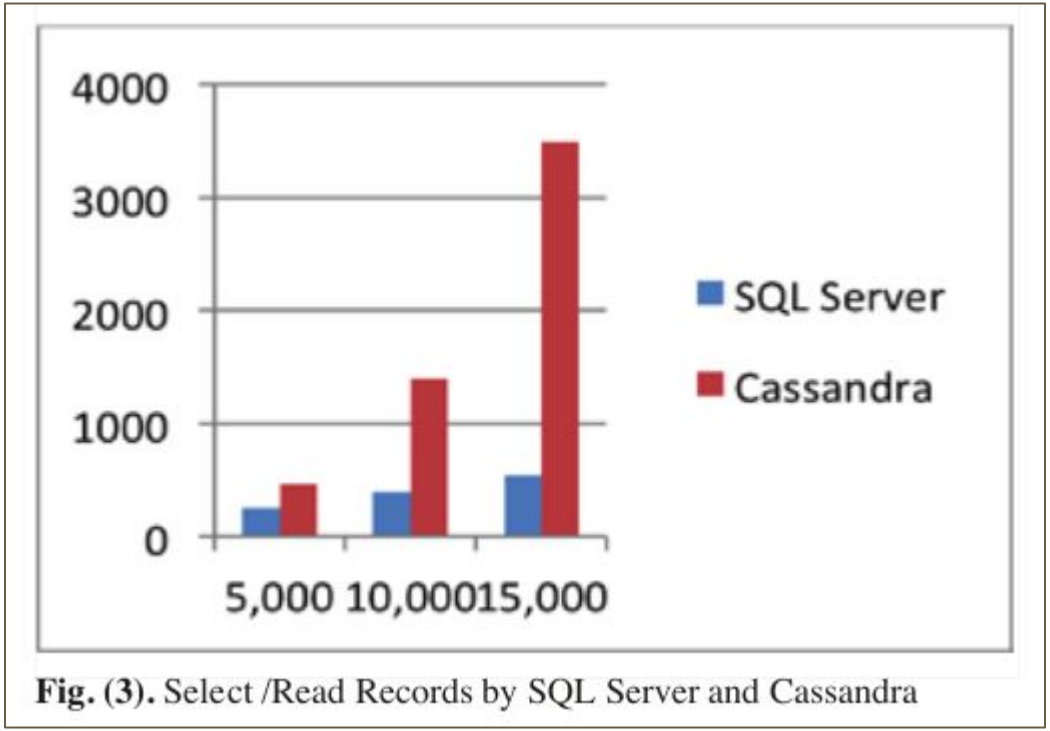

Fig. (3). Select /Read Records by SQL Server and Cassandra

Bonus Ablation Study (2 node vs 3 node)

- 3 node takes longer than 2 nodes
- However, difference cannot be said as significant
- Need more samples

Quality Attribute Analysis - Conclusion

Availability	**H**	Cassandra has higher availability than Postgres because of replication ability. If nodes are down in cluster, then other nodes have backup.
Performance	**H**	Postgres seems to be better in terms of read performance according to our testing.
Scalability	**H**	Cassandra will definitely be more scalable. Postgres only supports vertical scalability whereas Cassandra supports horizontal scalability. We can easily add more nodes whether in same local machine or on a remote machine. See README for info.

Conclusion and Analysis

- We couldn’t do more than 500k - took too long time
- Reason for inferior Cassandra Performance may be the lack of FKs and joins in query
- If we were having joins in our queries, then Cassandra would be superior.
- Perhaps indexing and in-memory databases may be better
- Possible limitations in the use of Docker containers
- Possible config issues

Future work:

- Youtube Integration will be added later as we wanted to focus on finding a suitable architecture first before continuing development

USE CASES
&
UI DESIGN

Use Case: SearchKeyword

Use case name	SearchKeyword
Participating actors	Initiated by **User** Communicates with **System**
Flow of events	1. The **User** types the keyword in the search bar and presses the search button. 2. The System searches for the keyword in the transcripts of all videos in the database. 3. The **System** displays the search results. 4. The **User** chooses the result that they desire. 5. The **User** plays the video.
Entry condition	● The **User** is logged in.
Exit condition	● The **User** finds the video that they are searching for and clicks on the video link to play it. ● The **User** cancels the search.

Use Case: SearchKeyword

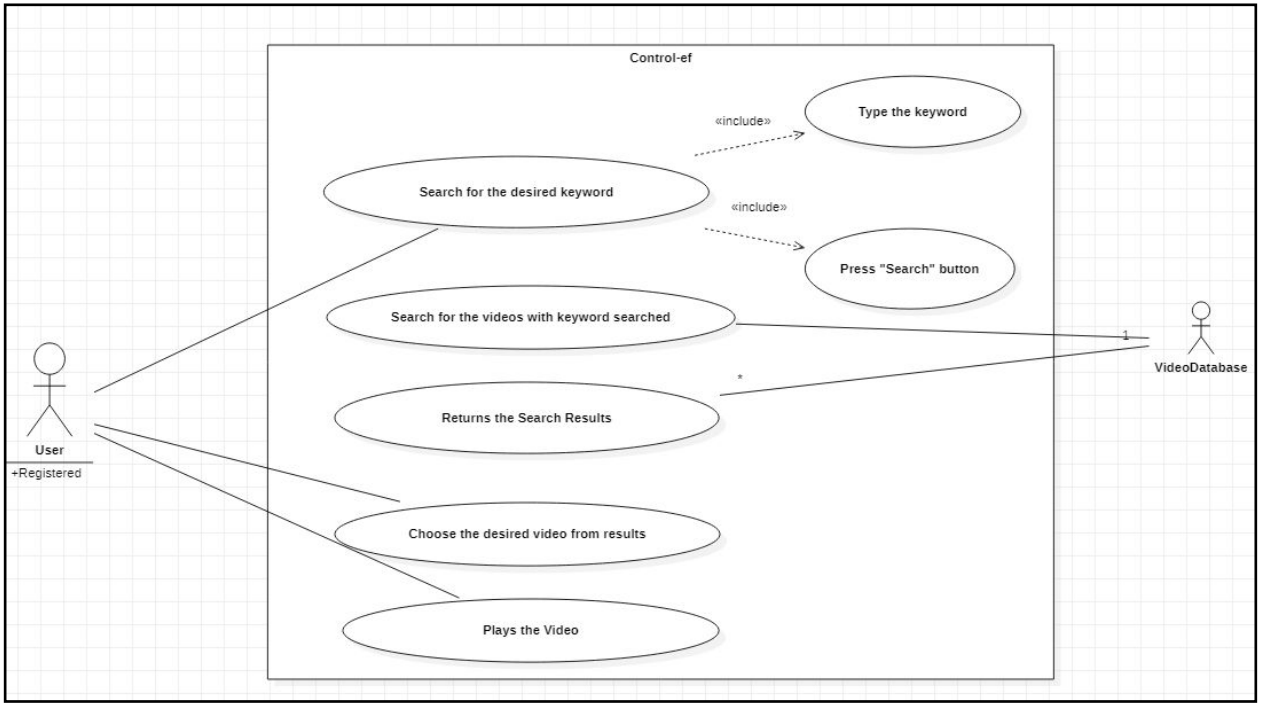

Use Case: SearchTag

Use case name	SearchTag
Participating actors	Initiated by **User** Communicates with **System**
Flow of events	1. The **User** types the tag in the search bar and presses the search button. 2. The **System** searches for all videos under that predefined tag in the database. 3. The **System** displays the search results. 4. The **User** chooses the result that they desire. 5. The **User** plays the video.
Entry condition	• The **User** is logged in.
Exit condition	• The **User** finds the video that they are searching for and clicks on the video link to play it. • The **User** cancels the search.

Use Case: SearchTag

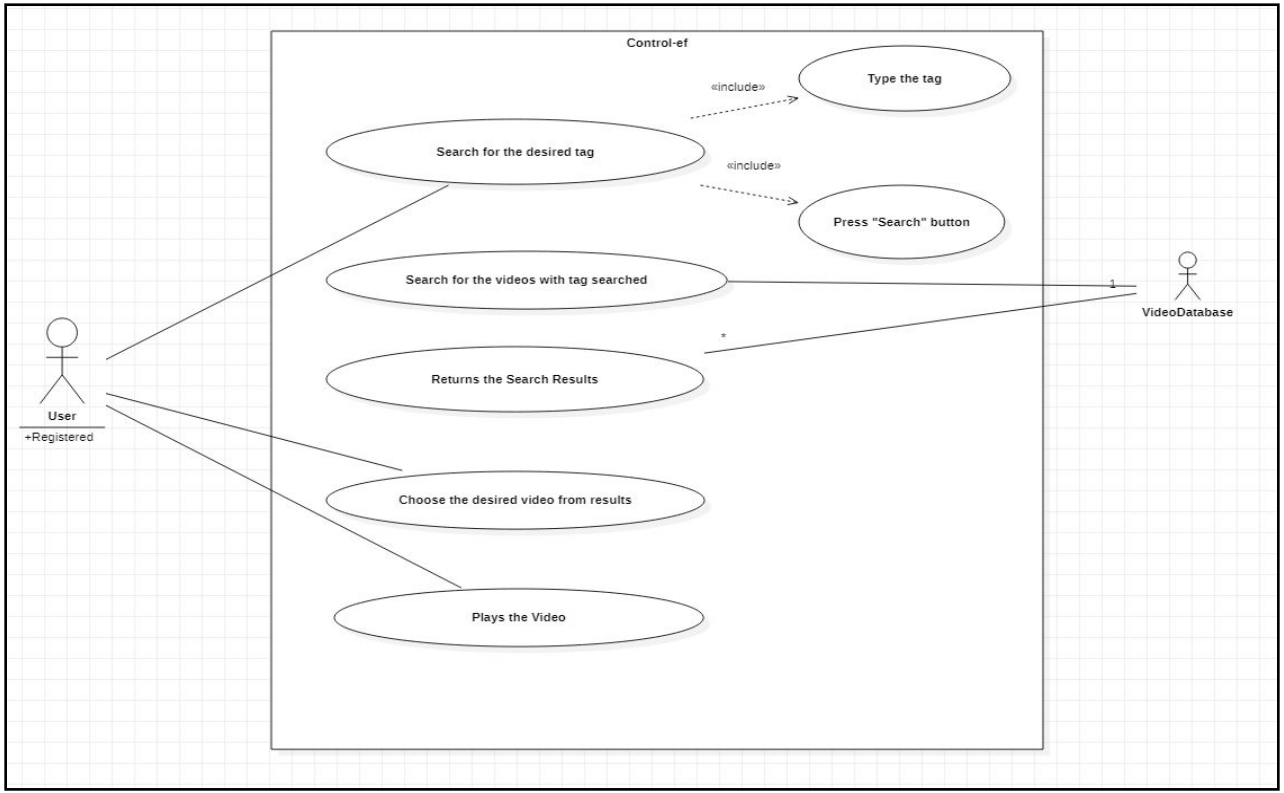

UI Design

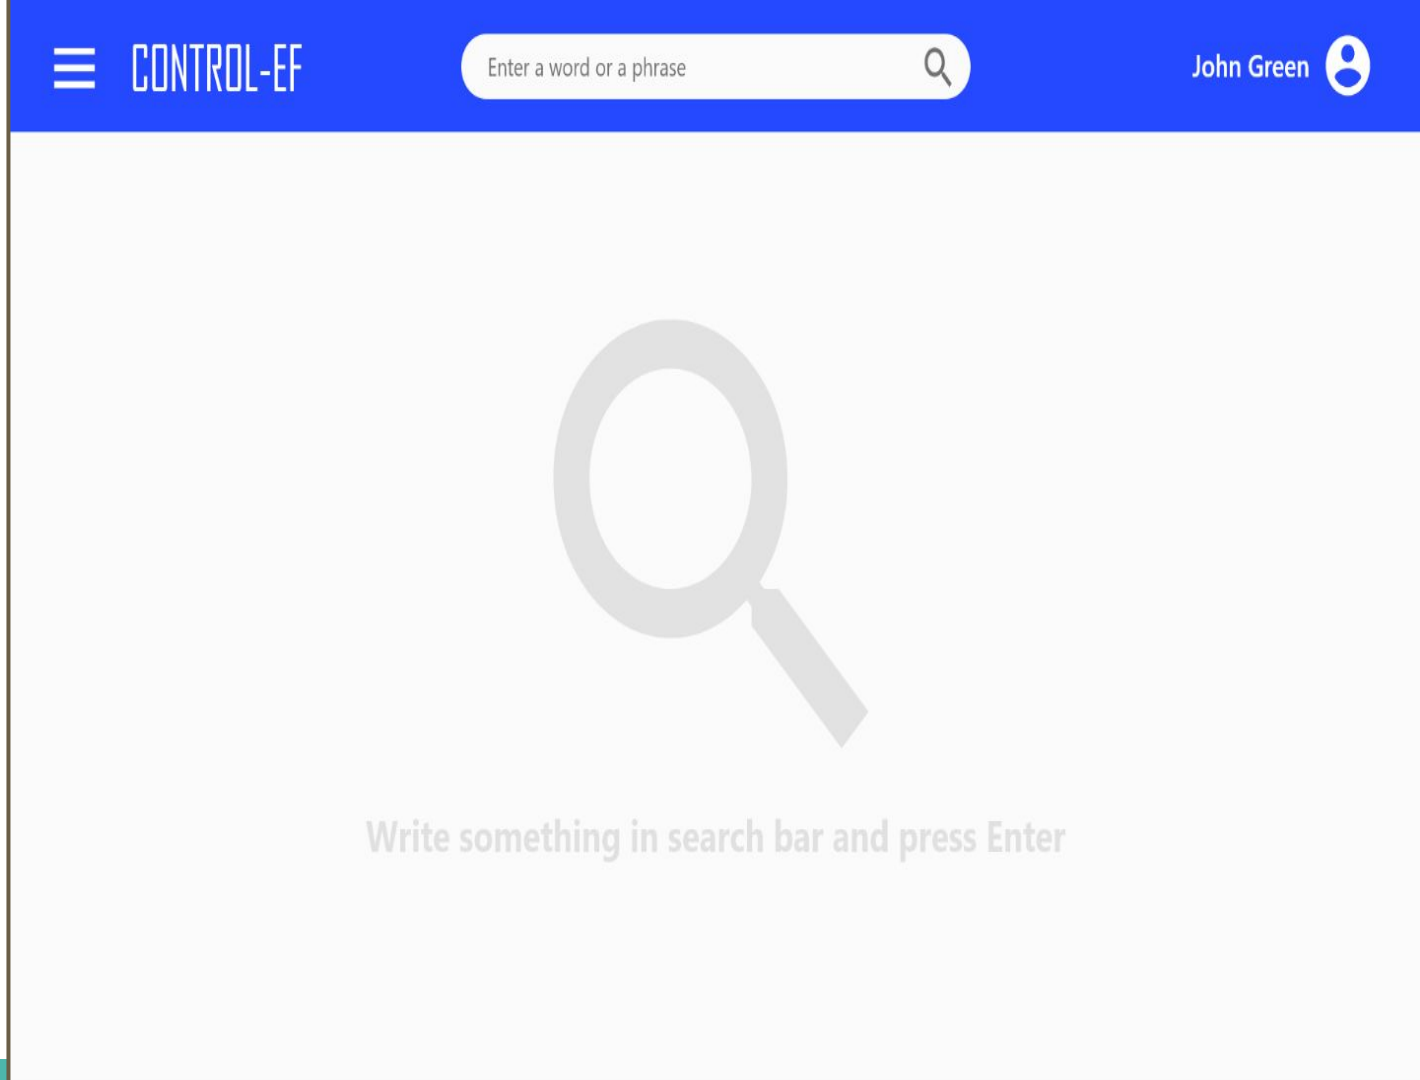
CONTROL-EF
Enter a word or a phrase
John Green
Write something in search bar and press Enter

UI Design

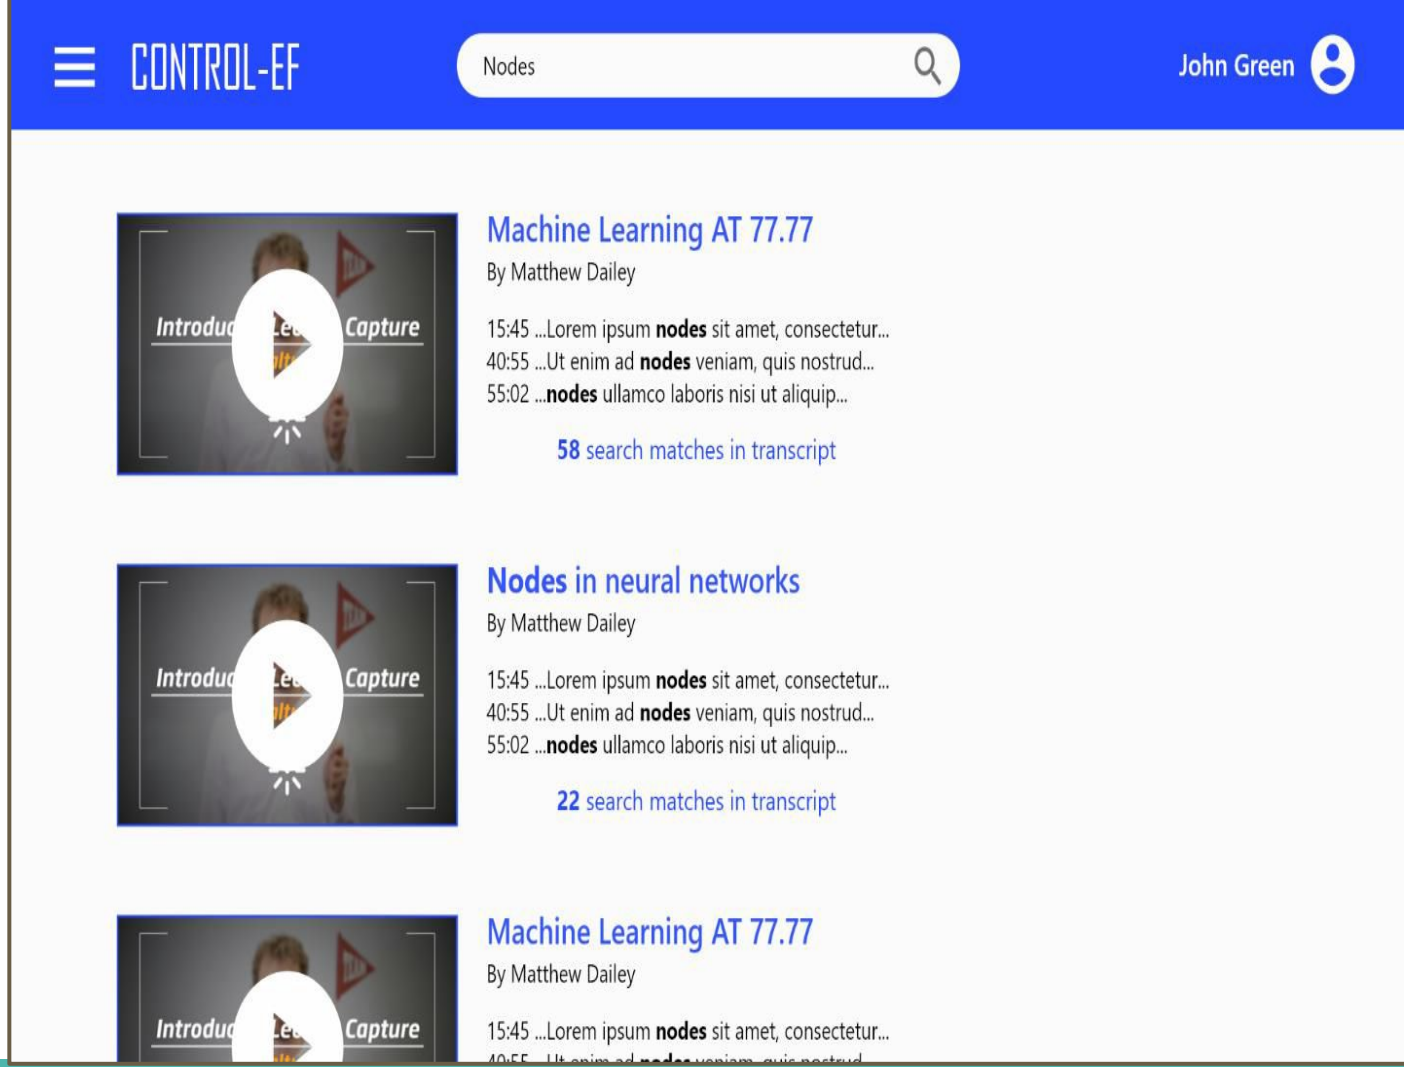
CONTROL-EF
Nodes
John Green
Machine Learning AT 77.77
By Matthew Dailey
15:45 ...Lorem ipsum nodes sit amet, consectetur...
40:55 ...Ut enim ad nodes veniam, quis nostrud...
55:02 ...nodes ullamco laboris nisi ut aliquip...
58 search matches in transcript
Nodes in neural networks
By Matthew Dailey
15:45 ...Lorem ipsum nodes sit amet, consectetur...
40:55 ...Ut enim ad nodes veniam, quis nostrud...
55:02 ...nodes ullamco laboris nisi ut aliquip...
22 search matches in transcript
Machine Learning AT 77.77
By Matthew Dailey
15:45 ...Lorem ipsum nodes sit amet, consectetur...

UI Design

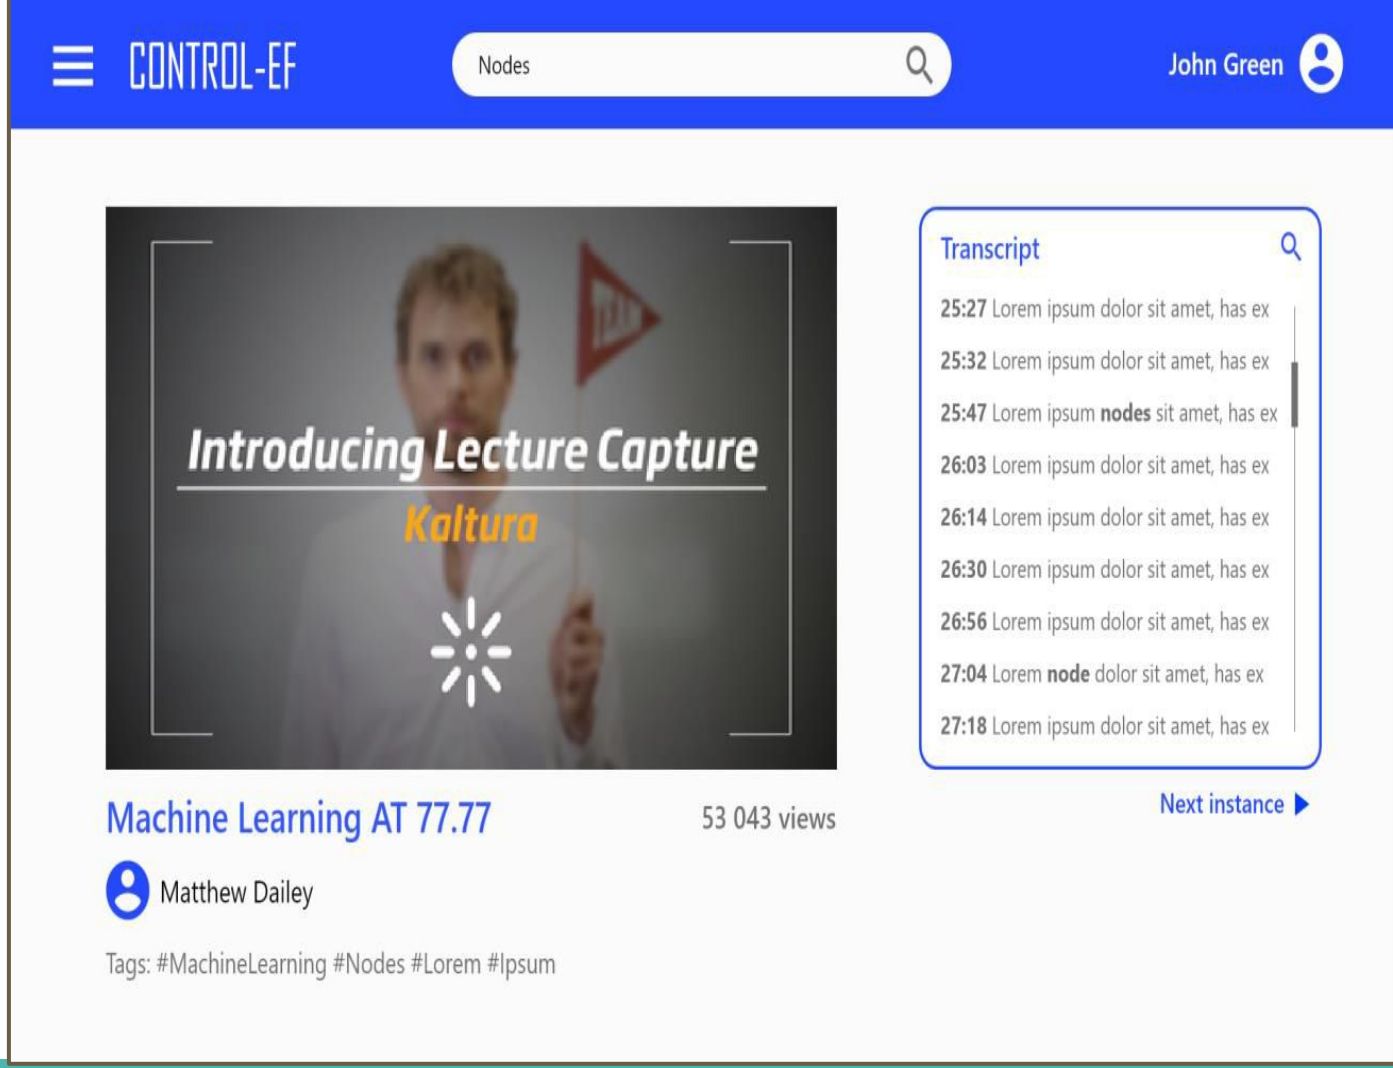
CONTROL-EF
Nodes
John Green
Introducing Lecture Capture
Kaltura
Transcript
25:27 Lorem ipsum dolor sit amet, has ex
25:32 Lorem ipsum dolor sit amet, has ex
25:47 Lorem ipsum nodes sit amet, has ex
26:03 Lorem ipsum dolor sit amet, has ex
26:14 Lorem ipsum dolor sit amet, has ex
26:30 Lorem ipsum dolor sit amet, has ex
26:56 Lorem ipsum dolor sit amet, has ex
27:04 Lorem node dolor sit amet, has ex
27:18 Lorem ipsum dolor sit amet, has ex
Next instance
Machine Learning AT 77.77
53 043 views
Matthew Dailey
Tags: #MachineLearning #Nodes #Lorem #Ipsum

Project Demo

Conclusion

- New problems always come up after coding
- Scope creep/inflation when we try something new
- Finding a suitable architecture or architectures
- Development/production environment discrepancy


thank you
danke
謝謝
ngiyabonga
teşekkür ederim
gracias
спасибо
obrigado
merci
dziękuję
bedankt
go raibh maith agat
mochchakkeram
sukriya
kop khun krap
terima kasih
grazie
arigatô
dakujem
tapadh leat
dank je
hvala
mauruuru
sagolun
takk
taiku
ευχαριστώ
감사합니다
ধন্যবাদ
рахмат
Баярлалаа
merci